

# H-DROID CLUSTER



Bruzzo Paolo and Casula Dario  
{paolo.bruzzo; dario.casula}@mail.polimi.it  
HPPS Project at Polimi  
March - June 2014

Tutor: Ferroni Matteo  
Professors: Donatella Sciuto, Marco D. Santambrogio

# CONTEXT DEFINITION

Nowadays... more mobile devices than humans<sup>1</sup>

Idle periods are quite common:  
(e.g. night time)



Potential huge amount of distributed  
computational power

How to exploit it ?

# PROJECT GOAL

to build an **heterogeneous cluster**:

distributed computation

which performance ?



previous attempt:

2012-droidCluster, homogeneous smartphone cluster<sup>2</sup>

limitation: all equal devices





# TECHNOLOGIES

Computer networks:

**MPICH2**, message passing interface  
devices communication

Measuring performance:

**HPL**, High Performance Linpack  
benchmark specifically written for clusters



# HPL ALGORITHM

Resolution of a linear algebraic system of N linear equations

Matrix divided into blocks

A <sub>11</sub>	A <sub>12</sub>	A <sub>13</sub>	A <sub>14</sub>	A <sub>15</sub>	A <sub>16</sub>	A <sub>17</sub>	A <sub>18</sub>
A <sub>21</sub>	A <sub>22</sub>	A <sub>23</sub>	A <sub>24</sub>	A <sub>25</sub>	A <sub>26</sub>	A <sub>27</sub>	A <sub>28</sub>
A <sub>31</sub>	A <sub>32</sub>	A <sub>33</sub>	A <sub>34</sub>	A <sub>35</sub>	A <sub>36</sub>	A <sub>37</sub>	A <sub>38</sub>
A <sub>41</sub>	A <sub>42</sub>	A <sub>43</sub>	A <sub>44</sub>	A <sub>45</sub>	A <sub>46</sub>	A <sub>47</sub>	A <sub>48</sub>
A <sub>51</sub>	A <sub>52</sub>	A <sub>53</sub>	A <sub>54</sub>	A <sub>55</sub>	A <sub>56</sub>	A <sub>57</sub>	A <sub>58</sub>
A <sub>61</sub>	A <sub>62</sub>	A <sub>63</sub>	A <sub>64</sub>	A <sub>65</sub>	A <sub>66</sub>	A <sub>67</sub>	A <sub>68</sub>
A <sub>71</sub>	A <sub>72</sub>	A <sub>73</sub>	A <sub>74</sub>	A <sub>75</sub>	A <sub>76</sub>	A <sub>77</sub>	A <sub>78</sub>
A <sub>81</sub>	A <sub>82</sub>	A <sub>83</sub>	A <sub>84</sub>	A <sub>85</sub>	A <sub>86</sub>	A <sub>87</sub>	A <sub>88</sub>

Grid of processes

A <sub>11</sub>	A <sub>14</sub>	A <sub>17</sub>	A <sub>12</sub>	A <sub>15</sub>	A <sub>18</sub>	A <sub>13</sub>	A <sub>16</sub>
A <sub>31</sub>	A <sub>34</sub>	A <sub>37</sub>	A <sub>32</sub>	A <sub>35</sub>	A <sub>38</sub>	A <sub>33</sub>	A <sub>36</sub>
A <sub>51</sub>	A <sub>54</sub>	A <sub>57</sub>	A <sub>52</sub>	A <sub>55</sub>	A <sub>58</sub>	A <sub>53</sub>	A <sub>56</sub>
A <sub>71</sub>	A <sub>74</sub>	A <sub>77</sub>	A <sub>72</sub>	A <sub>75</sub>	A <sub>78</sub>	A <sub>73</sub>	A <sub>76</sub>
A <sub>21</sub>	A <sub>24</sub>	A <sub>27</sub>	A <sub>22</sub>	A <sub>25</sub>	A <sub>28</sub>	A <sub>23</sub>	A <sub>26</sub>
A <sub>41</sub>	A <sub>44</sub>	A <sub>47</sub>	A <sub>42</sub>	A <sub>45</sub>	A <sub>48</sub>	A <sub>43</sub>	A <sub>46</sub>
A <sub>61</sub>	A <sub>64</sub>	A <sub>67</sub>	A <sub>62</sub>	A <sub>65</sub>	A <sub>68</sub>	A <sub>63</sub>	A <sub>66</sub>
A <sub>81</sub>	A <sub>84</sub>	A <sub>87</sub>	A <sub>82</sub>	A <sub>85</sub>	A <sub>88</sub>	A <sub>83</sub>	A <sub>86</sub>

Matrix order and block size choices

# WHY ANDROID

How to run MPI and HPL on Android ?



**Debian on Android !**



Chroot into Debian: same kernel of Android




# INTO DEBIAN

Bash script to chroot into Debian



`root@phone# sh boot.sh`



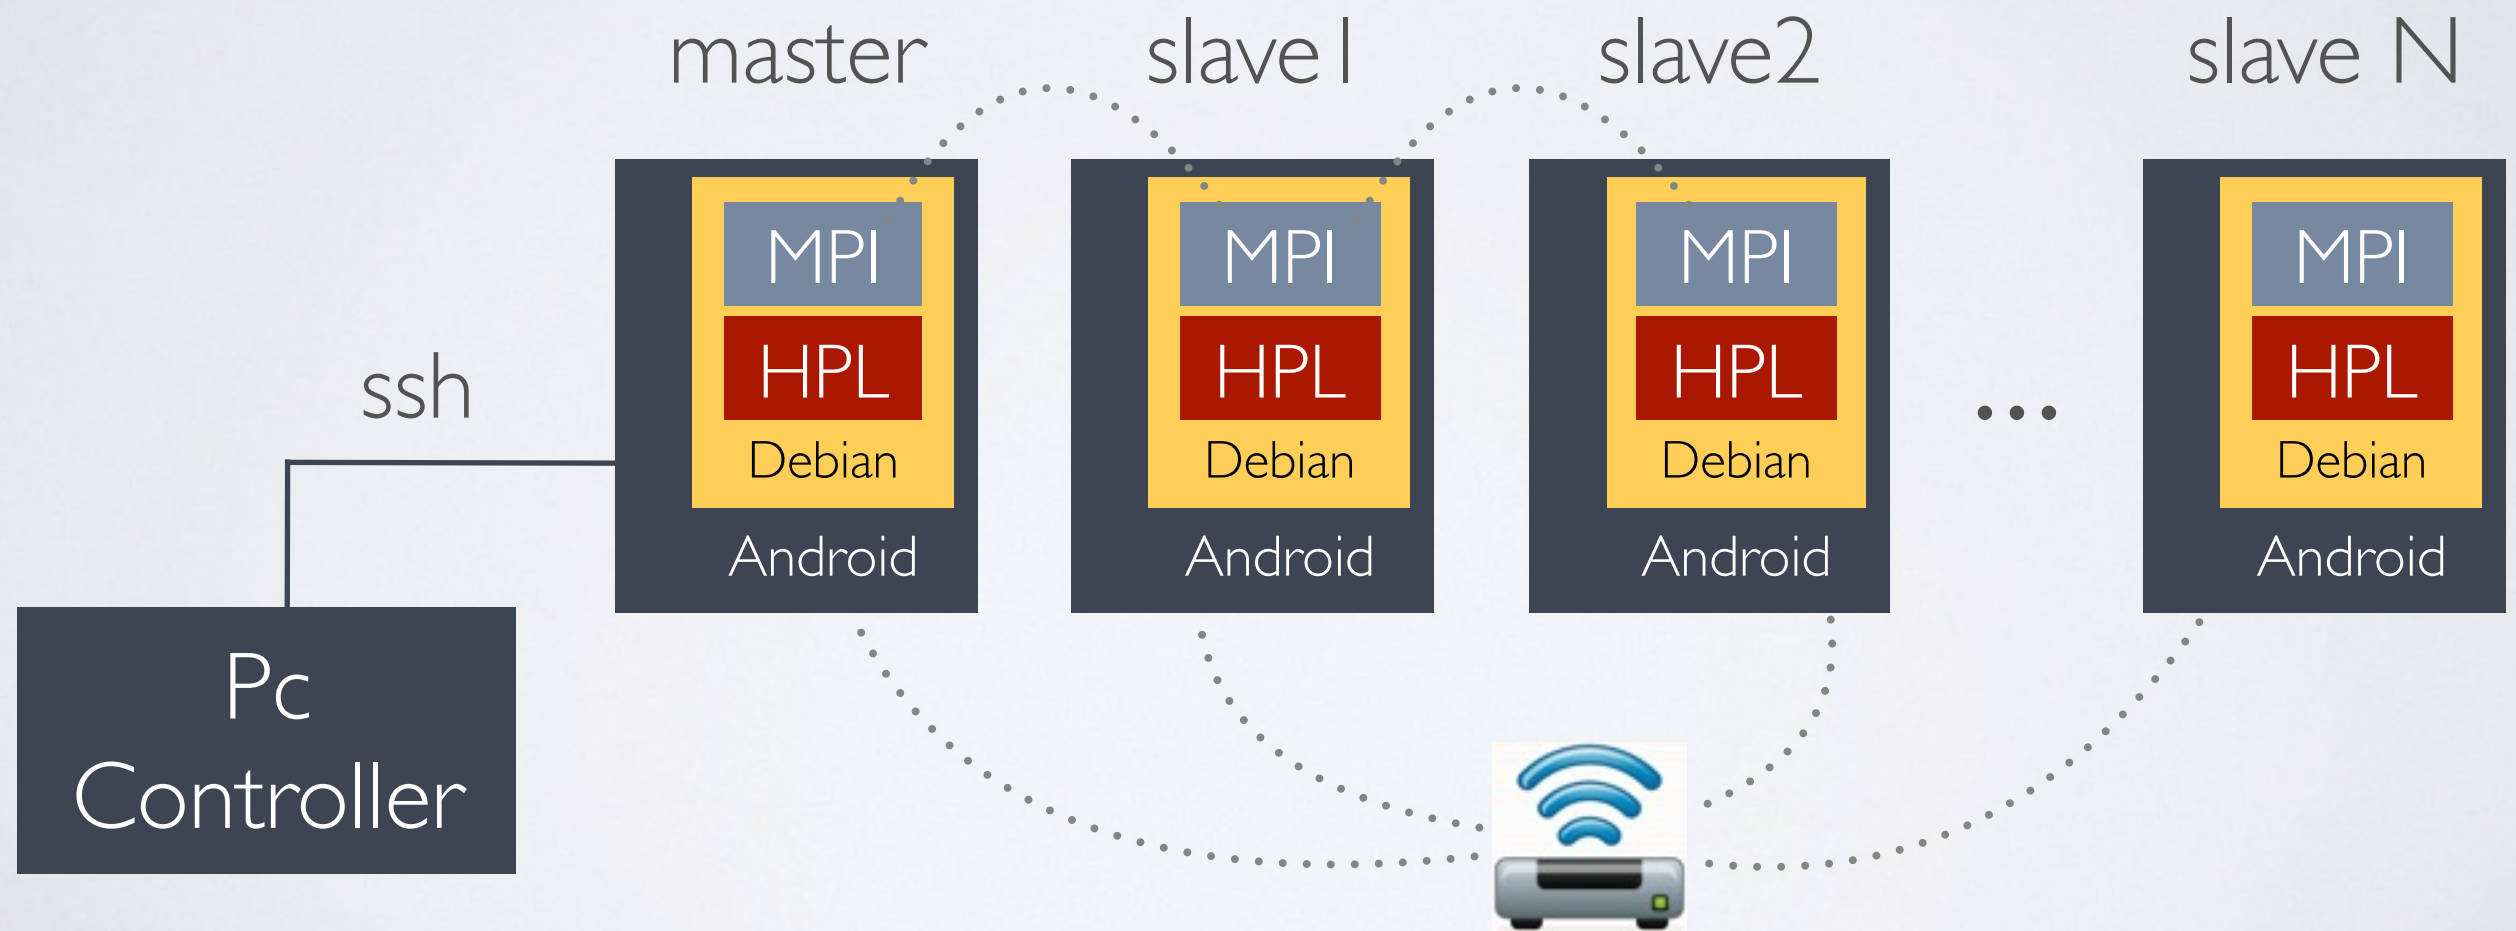
C program to launch the MPI cluster ring

How to tune the benchmark ?

# CASE STUDY

Heterogeneous cluster of Android Devices :


- Master distributes workloads
- Wireless connection





# TYPES OF CLUSTERS

Homogeneous  
all equal devices




Free memory

OS memory

CPU

Heterogeneous  
our case




Trivial benchmark tuning

Hard benchmark tuning

Choice of MATRIX ORDER ?

# TUNING: MATRIX ORDER


Samsung S5570 - Trade Off: Time and Performance



# TUNING: BLOCK SIZE


No precise rules, test different measures

Devices return same performance with different sizes





# CLUSTER PERFORMANCE



# PROCESSES ALLOCATION

Some devices are more powerful, why not to give them more work to do?

A <sub>11</sub>	A <sub>14</sub>	A <sub>17</sub>
A <sub>31</sub>	A <sub>34</sub>	A <sub>37</sub>
A <sub>51</sub>	A <sub>54</sub>	A <sub>57</sub>
A <sub>71</sub>	A <sub>74</sub>	A <sub>77</sub>
A <sub>21</sub>	A <sub>24</sub>	A <sub>27</sub>
A <sub>41</sub>	A <sub>44</sub>	A <sub>47</sub>
A <sub>61</sub>	A <sub>64</sub>	A <sub>67</sub>
A <sub>81</sub>	A <sub>84</sub>	A <sub>87</sub>



dual core


A <sub>22</sub>	A <sub>25</sub>	A <sub>28</sub>
A <sub>42</sub>	A <sub>45</sub>	A <sub>48</sub>
A <sub>62</sub>	A <sub>65</sub>	A <sub>68</sub>
A <sub>82</sub>	A <sub>85</sub>	A <sub>88</sub>



single core


# 1 PROCESS PER CORE

■ Device Local Performance ● Theoretical ● Real ● Previous Real






# THE BOTTLENECK



# FUTURE WORK

What happens with tens /  
hundreds of devices ?



QUESTIONS ?



# REFERENCES

1. [digitaltrends.com](http://digitaltrends.com): Number of mobile phones to exceed world population by 2014
2. Büsching, Schildt, Wolf: "DroidCluster: Towards Smartphone Cluster Computing"; 2012, Technische Universität Braunschweig
3. J. Dongarra, P. Luszczek, A. Petitet, "The linpack benchmark: Past, present and future." 2003
4. Daniel Loreto, Erik Nordlander, Adam Oliner, "Benchmarking a Large-Scale Heterogeneous Cluster", MIT, 2005
5. [netlib.org: /benchmark/hpl/tuning.html](http://netlib.org/benchmark/hpl/tuning.html)